Vikrant Choudhary
Email: vikrant.choudhary@mailbox.in | Phone: +91 95075 93022 | Location: Gurugram, India
LinkedIn: https://linkedin.com/in/vikrant-choudhary | GitHub: https://github.com/vchoudhary

PROFESSIONAL SUMMARY

Senior Machine Learning Engineer with 5.5 years of production experience across e-commerce and mobility. Deep hands-on expertise in PyTorch, TensorFlow, FastAPI, Docker. Known for measurable impact on latency, reliability and cost, and for mentoring engineers.

TECHNICAL SKILLS

PyTorch, TensorFlow, FastAPI, Docker, AWS, Hugging Face, MLOps, Scikit-Learn, SQL, Python, Pandas, NumPy, Airflow, CI/CD, Machine Learning

EDUCATION

B.Tech in Computer Science and Engineering
Vellore Institute of Technology | CGPA: 8.4/10 | Graduated: 2021

WORK EXPERIENCE

Senior Machine Learning Engineer | Basalt Mobility (Oct 2023 - Present)
• Designed a feature store on PostgreSQL and Redis with point-in-time correctness; online feature reads at 3 ms p99.
• Reduced AWS GPU spend by 36% (about $31k per month) through INT8 quantization, dynamic batching and spot-instance training.
• Implemented drift and data-quality monitoring with Prometheus and Evidently that caught 11 silent regressions before reaching 40,000 daily users.
• Deployed RoBERTa ONNX models with FastAPI and Triton serving 80k daily requests at 45 ms p95 latency on AWS ECS.

Machine Learning Engineer | Pelican Health Data (Apr 2021 - Oct 2023)
• Cut training time for a 340M-parameter PyTorch model by 52% using mixed precision, gradient checkpointing and DDP on 16 GPUs.
• Shipped a TensorFlow Serving to PyTorch migration for 14 legacy models with zero downtime and a 33% latency improvement.
• Led a squad of 5 engineers, introduced PyTest-based model contracts and lifted coverage of the training stack to 84%.

KEY PROJECTS

Edge Anomaly Detector
• LSTM autoencoder in PyTorch quantised to INT8 and deployed to ARM devices; 95% precision on 400,000 sensor windows.

Semantic Resume Search
• RAG pipeline using sentence-transformers, FAISS and a FastAPI backend; returns results in under 62 ms across 40,000 documents.

CERTIFICATIONS & AWARDS

• NVIDIA DLI: Fundamentals of Deep Learning

- Runner-up, Google Cloud Hackathon India 2024